



Principles of Computer System Design (PCSD)

Marcos Vaz Salles
Assistant Professor, DIKU

Matchmaking event: Data and Biotech

What:

- Meet the following companies: **Novozymes**, **Novo Nordisk**, **Lundbeck**, **Biomediq**, **Chr. Hansen**, **CVIVA** and **Statens Serum Institut**.
- Each company will present problems/challenges that they want **YOU** to solve in a **Master Thesis**.
- After the presentations you will have the opportunity to meet up with 2-3 companies for a one-on-one discussion about future collaboration.

When:

November 22nd, 1pm-6pm – **THIS FRIDAY**
BUT REGISTER BY TOMORROW

Where:

HCØ, Universitetsparken 5

Sign up and read more: <http://www.science.ku.dk/matchmaking/>



Why study computer systems?

- The IBM/Microsoft/Oracle question


How can I program large systems with clean interfaces and high performance?

- The Amazon/Facebook/Google question

How can I understand the guarantees and reliability of scalable services offered to me on the cloud?

- The Cloudera/Greenplum/Teradata question

How do I build systems to process TBs to PBs of data?




Intro Video by Michael Brodie

Computer Science 2.0
presented at VLDB 2007, Vienna, Austria



What should we learn in this course?



- Knowledge
 - Describe the design of transactional and distributed systems
 - Explain how to enforce modularity through a client-service abstraction
 - Explain techniques for large-scale data processing
- Skills
 - Implement systems that include mechanisms for modularity, atomicity, and fault tolerance
 - Structure and conduct experiments to evaluate a system's performance
- Competences
 - Discuss design alternatives for a computer system, identifying system properties as well as mechanisms for improving performance
 - Analyze protocols for concurrency control and recovery, as well as for distribution and replication.
 - Apply principles of large-scale data processing to concrete problems.



PCSD: What will we study?

- **Fundamentals**


- Abstractions: interpreters, memory, communication links
- Modularity with clients and services, RPC
- Techniques for performance, e.g., concurrency, fast paths, dallying, batching, speculation



Property: Strong Modularity

PCSD: What will we study?

- **Concurrency Control and Recovery**
 - Two-phase locking
 - Serializability, schedules
 - Optimistic and multi-version approaches to concurrency control
 - Recovery concepts
 - ARIES recovery algorithm



Properties: Atomicity and Durability

PCSD: What will we study?

- **Communication**
 - Message queues, streams, multicast, BASE
 - End-to-end argument
- **Reliability & Distribution**
 - Reliability concepts
 - Replication techniques
 - Topics in coordination and distributed transactions



Property: High Availability

PCSD: What will we study?

- **Data Processing**
 - Operators
 - External sorting
 - Hash- and sort-based techniques for multiple operations (e.g., duplicate elimination & grouping, set operations, joins)
 - Parallelism



**Property: Scalability with
Data Size**


PCSD: What will we study?

- **Experimental Design**
 - Performance metrics, workloads
 - Structuring and conducting simple experiments



PCSD: What will we study?

- Guest Lectures
 - **Date: December 17**
 - **Service-Oriented Architecture (SOA) in the Real World**, by Morten Steffensen of **Netcompany**
 - **Recovery in Practice**, by Paz Padilla Thygesen of **IBM**



References & Course Materials

- Course webpage
 - Kurser:
<http://kurser.ku.dk/course/ndaa09004u/2013-2014>
- Course materials in Absalon
 - Tentative course syllabus
 - Includes readings for after each class
 - Slides before each class
 - Assignments & Feedback
 - Message forums
- Please always post your questions in Absalon
 - Your colleagues can profit too!



References & Course Materials

- Book
 - Principles of Computer System Design (PCSD): DIKU Course Compendium. Collected references from sources cited therein, organized by Marcos Vaz Salles and Michael Kirkedal Carøe.
- Papers & other references
 - Vast majority listed in the syllabus
 - A few more will come as we go
 - Optional references for more depth



Team

- Lecturers
 - Marcos Vaz Salles
 - vmarcos@diku.dk
 - Jyrki Katajainen
 - jyrki@diku.dk
 - Office hours:
 - By email appointment
- TAs
 - Vivek Shah
 - Ashwini Satish Joshi
 - Jacob Jepsen
 - Meet them in the TA sessions on Thursdays!



Weekly Schedule

- Lectures Tuesdays and Thursdays, 10am-12pm
 - Two 45 min sessions, 15 min break, with lecturer
 - Participation will be encouraged 😊
- TA sessions Thursdays 1pm-3pm
 - Session length according to need
 - TAs will guide most of those
 - Exercises
 - Assignment work time and Q&A

Learning is the main goal!



First Steps

- **Java Warm-up Exercise**
 - Available on Absalon
 - If you passed Advanced Java, you do not need it 😊
 - If you did not take Advanced Java, the warm-up assignment will tell you the level of Java you **need** for this course
- **First TA Session**
 - This Thursday, **November 21**
 - **Brief review** of Java Warm-up Exercise
 - Setup of **optional Windows Azure** cloud service
 - Generous gift by Microsoft
 - Allows you to learn while using a cloud service
 - We have been awarded one pass per student
 - Roughly two small instances for 5 months



Assignments

- **4 + 1 take-home assignments**
 - **Groups: 2 people** strongly recommended
 - Each assignment worth **10 points**
 - Minimum of **30 points** to qualify for exam
- **First 4 assignments**
 - Build **specific** skills and concepts on **weekly** basis
 - Include both theory exercises and programming
 - **Due dates: December 1, December 8, December 15, December 22**
- **Final 5th assignment**
 - **Integrates** multiple skills and concepts into a single assignment
 - **Exam-style:** Based on last year's exam
 - Includes both theory exercises and programming
 - **Due date: January 8**



Exam

- **Exam format**

- 5-day take-home assignment with external grading on 7-point scale, between **January 15** and **January 21**
- Must be solved individually, no groups allowed
- Includes both theory exercises and programming
- Similar in structure to Assignment 5
- Submission in Absalon

Academic Integrity
taken very seriously



Acknowledgements

- Many of the slides in this course are based on or reproduce material kindly made available by Jerome Saltzer & M. Frans Kaashoek & Robert Morris (MIT, PCSD textbook material), Johannes Gehrke (Cornell, Ramakrishnan & Gehrke textbook), Gustavo Alonso (ETH Zurich, EAI course), Nesime Tatbul (ETH Zurich), James Kurose & Keith Ross (U Mass Amherst & NYU, networking textbook), Jens Dittrich (Saarland University)




PCSD: Evaluation and Evolution

- PCSD is evolving based on your feedback
 - Positive sentiment last year after changes to syllabus
 - Feedback from last year
 - Syllabus is great!
 - Assignments were too steep in workload / learning curve
- We are listening
 - Course syllabus kept this year
 - Assignments broken down into 4 + 1 model
 - Skills building course: **Advanced Java**
- Helping out
 - Come and talk to me, give us your feedback!
 - Make sure to fill out our evaluation questionnaires – WE READ THEM



Questions so far?






What should we learn today?

- Identify the fundamental abstractions in computer systems and their APIs, including memory, interpreters, communication links
- Explain how names are used in the fundamental abstractions
- Design a top-level abstraction, respecting its correspondent API, based on lower-level abstractions
- Discuss performance and fault-tolerance aspects of such a design




The Central Trade-off: Abstractions, Performance, Fault-Tolerance



Fundamental abstractions

- **Memory**

- Read/write




- **Interpreters**

- Instruction repertoire
- Environment
- Instruction pointer

(loop (print (eval (read))))

Names
make
connections



- **Communication links**

- Send/receive



Source: Saltzer & Kaashoek & Morris (partial)

Examples of Names


- R1
- 1742
- 18.7.22.69
- web.mit.edu
- http://web.mit.edu/6.033
- 6.033-staff@mit.edu
- amsterdam
- /mit/6.033/www
- foo.c
- .. (as in cd .. or ls ..)
- wc
- (617)253-7149, x37149
- 021-84-2030

address is overloaded
name with location info
(e.g., LOAD 1742, R1)

Names require a
mapping scheme



Name Mapping




- How can we map names?
- Table lookup
 - Files inside directories
- Recursive lookup
 - Path names in file systems or URLs
- Multiple lookup
 - Java class loading



Fundamental abstractions

- **Memory**


- Read/write



- **Interpreters**

- Instruction repertoire
- Environment
- Instruction pointer

`(loop (print (eval (read))))`



- **Communication links**

- Send/receive



Source: Saltzer & Kaashoek & Morris (partial)

Memory

- **Memory**

- `READ(name) → value`
- `WRITE(name, value)`




- Examples of Memory

- Physical memory (RAM)
- Multi-level memory hierarchy (registers, caches, RAM, flash, disk, tape)
- Address spaces and virtual memory with paging
- Transactional memory (hardware and software variants)
- Database storage engines
- Key-value stores (e.g., Cassandra, Dynamo)



How would you design a two-level memory abstraction consolidating disk and RAM?

- Characteristics of storage technologies
 - **RAM**: high cost per gigabyte, low latency, volatile
 - **Disk**: low cost per gigabyte, high latency, nonvolatile
- Design top-level abstraction respecting *Memory API*
- Abstraction must:
 - Address as much data as fits in disk
 - Use fixed-size pages for disk transfers
 - Use RAM efficiently to provide for low latency (on average)
 - Neither disk nor memory directly exposed, only READ/WRITE API



Write the pseudocode down!



Address Space Mapping

- Address spaces modular way to multiplex memory
- Naming scheme translating virtual into physical addresses
- Page map
 - Updated by kernel code
 - Lookup implemented in hardware
 - Concerns: Protection (Pr), representation, efficiency

Page Map

Page#	Block#	Pr
42	2	RX
53	45	R
64	97	RW
...




Address Space Mapping: Introducing Disks

- Use disk to store more blocks
- Pages may be either in memory or on disk
- Resident bit (R)
 - Access to non-resident pages results in page faults
- Page Fault
 - An indirection exception for missing pages

Page Map

Page#	Block#	Pr	R	D	P
42	2	RX	1	1	0
53	45	R	1	0	1
64	D-42	RW	0	0	0
...




Address Space Mapping: Introducing Disks

- Handling page faults
 - Trap to OS handler
 - Handler loads block from disk and updates mapping
 - If memory full, must choose some *victim* block for replacement
 - Page replacement algorithm, e.g., LRU
- Other metadata
 - *Dirty bit (D)*: Only write page back when it has changed!
 - *Pin bit (P)*: do not remove certain pages (e.g., code of OS handler itself)

Page Map

Page#	Block#	Pr	R	D	P
42	2	RX	1	1	0
53	45	R	1	0	1
64	97	RW	1	0	0
...




Virtual Memory with Paging: Abstractions, Performance, Fault-Tolerance

- **Abstraction:** Do we have any guarantees on two concurrent threads writing to the same memory?
- **Performance:** Do we get average latency close to RAM latency?
- **Fault-Tolerance:** What happens on failure? Do we have any guarantees about the state that is on disk?

Page Map

Page#	Block#	Pr	R	D	P
42	2	RX	1	1	0
53	45	R	1	0	1
64	97	RW	1	0	0
...




Interpreters

- ***Interpreter***

- Instruction repertoire
- Environment
- Instruction pointer

(loop (print (eval (read))))



```
procedure INTERPRET()  
  do forever  
    instruction  $\leftarrow$  READ(instruction_pointer)  
    perform instruction in environment context  
    if interrupt_signal = TRUE then  
      instruction_pointer  $\leftarrow$  entry of INTERRUPT_HANDLER  
      environment  $\leftarrow$  environment of INTERRUPT_HANDLER
```

- Examples of Interpreters

- Processors (CPU)
- Programming language interpreters
- Frameworks, e.g., MapReduce
- Your own (layered) programs! (RPCs)

Source:
Saltzer &
Kaashoek &
Morris
(partial)



Communication links

- ***Communication links***

- `SEND(linkName, outgoingMessageBuffer)`
- `RECEIVE(linkName, incomingMessageBuffer)`



- **Examples of Communication Links**

- Ethernet interface
- IP datagram service
- TCP sockets
- Message-Oriented Middleware (MOM)
- Streams
- Multicast (e.g., CATOCS: Causal and Totally-Ordered Communication System)



Memory, Interpreters, Communication Links: Is that all there is?

- Other abstractions also useful!
- ***Synchronization***
 - Locks
 - Condition variables & monitors
(see, e.g., Chubby lock service from Google)
- ***Data processing***
 - Data transformations
 - Operators
(see, e.g., parallel implementations of SQL)



What should we learn today?



- Identify the fundamental abstractions in computer systems and their APIs, including memory, interpreters, communication links
- Explain how names are used in the fundamental abstractions
- Design a top-level abstraction, respecting its correspondent API, based on lower-level abstractions
- Discuss performance and fault-tolerance aspects of such a design

