



Modularity through Clients and Services, RPC Techniques for Performance


PCSD, Marcos Vaz Salles

Do-it-yourself Recap: Fundamental abstractions

- Which were the three fundamental abstractions?
- What were their APIs?
- Must these abstractions be implemented in a single node or can they be distributed? Give an example!

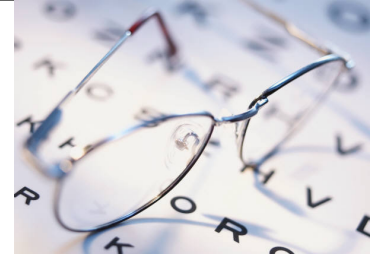


`(loop (print (eval (read))))`



Source: Saltzer & Kaashoek & Morris (partial)

What should we learn today?




- Recognize and explain modular designs with clients and services
- Predict the functioning of service calls under different RPC semantics and failure modes
- Identify different mechanisms to achieve RPCs
- Implement RPC services with an appropriate mechanism, such as web services
- Explain performance metrics such as latency, throughput, overhead, utilization, capacity, and scalability
- List common hardware parameters that affect performance
- Apply performance improvement techniques, such as concurrency, batching, dallying, and fast-path coding



Interpreters

- Interpreter
 - Instruction repertoire
 - Environment
 - Instruction pointer

(loop (print (eval (read))))



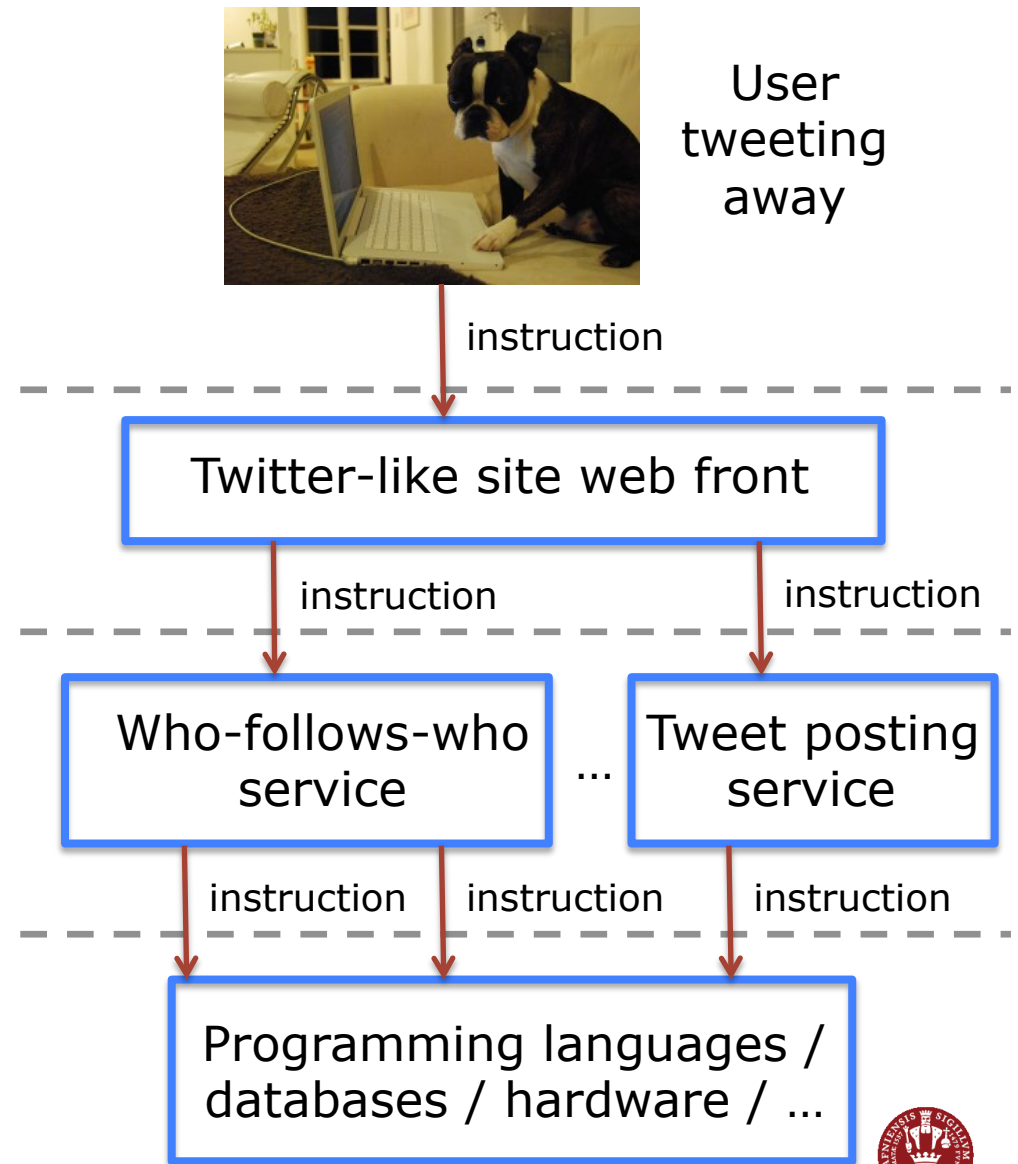
```
procedure INTERPRET()  
  do forever  
    instruction ← READ(instruction_pointer)  
    perform instruction in environment context  
    if interrupt_signal = TRUE then  
      instruction_pointer ← entry of INTERRUPT_HANDLER  
      environment ← environment of INTERRUPT_HANDLER
```

Your own programs often implement
specialized interpreters!



Layers and Modules

- Interpreters often organized in layers
- Modules
 - *Book glossary:* Components that can be separately designed / implemented / managed / replaced
 - “Instructions” of higher-level interpreters
 - Recursive: Can be whole interpreters themselves!





Cloud-Power @ Yahoo!

My Yahoo! | May 12, 2009 | Sign In | New here? Sign Up

Web Images Video Local Shopping More

Web Search

MY FAVORITES + Add

- View Yahoo! Sites
- Yahoo! Mail
- Autos
- eBay
- Finance (Down)
- Flickr
- Games
- Messenger
- Movies
- Music
- MySpace
- omg!
- Personals
- Sports
- TV
- Weather

TOP SEARCHES

1. ...
2. ...
3. ...
4. ...
5. ...
6. Leonard Nimoy
7. Fawcett
8. Grizzly Bear Attack
9. Stamps
10. TV Recaps

Search Index

Yahoo! Mail

Check your mail

Sign in

Don't have a Yahoo! ID? Sign up

Stay connected

Preview and keep up-to-date with your mail.

Attachment Storage

Image/Video Storage

TOYOTA

ROUGH JUNE 1ST.

See local offers

Toyota Sales Event - Ad Feedback

« Prev Next »

NEWS

A baby black jaguar is carried by its mom

What happens when modules fail with (unintended) errors?



Source: Raghu Ramakrishnan, LADIS 2009 (partial)



Cloud-Power @ Yahoo!



What happens when modules fail with
(unintended) errors?

Source: Raghu Ramakrishnan, LADIS 2009 (partial)



Isolating Errors: Enforced Modularity

Clients & Services

- Restrict communication to *messages only*
- Client request / Service response (or reply)
- Conceptually client and service in different computers

OS Virtualization

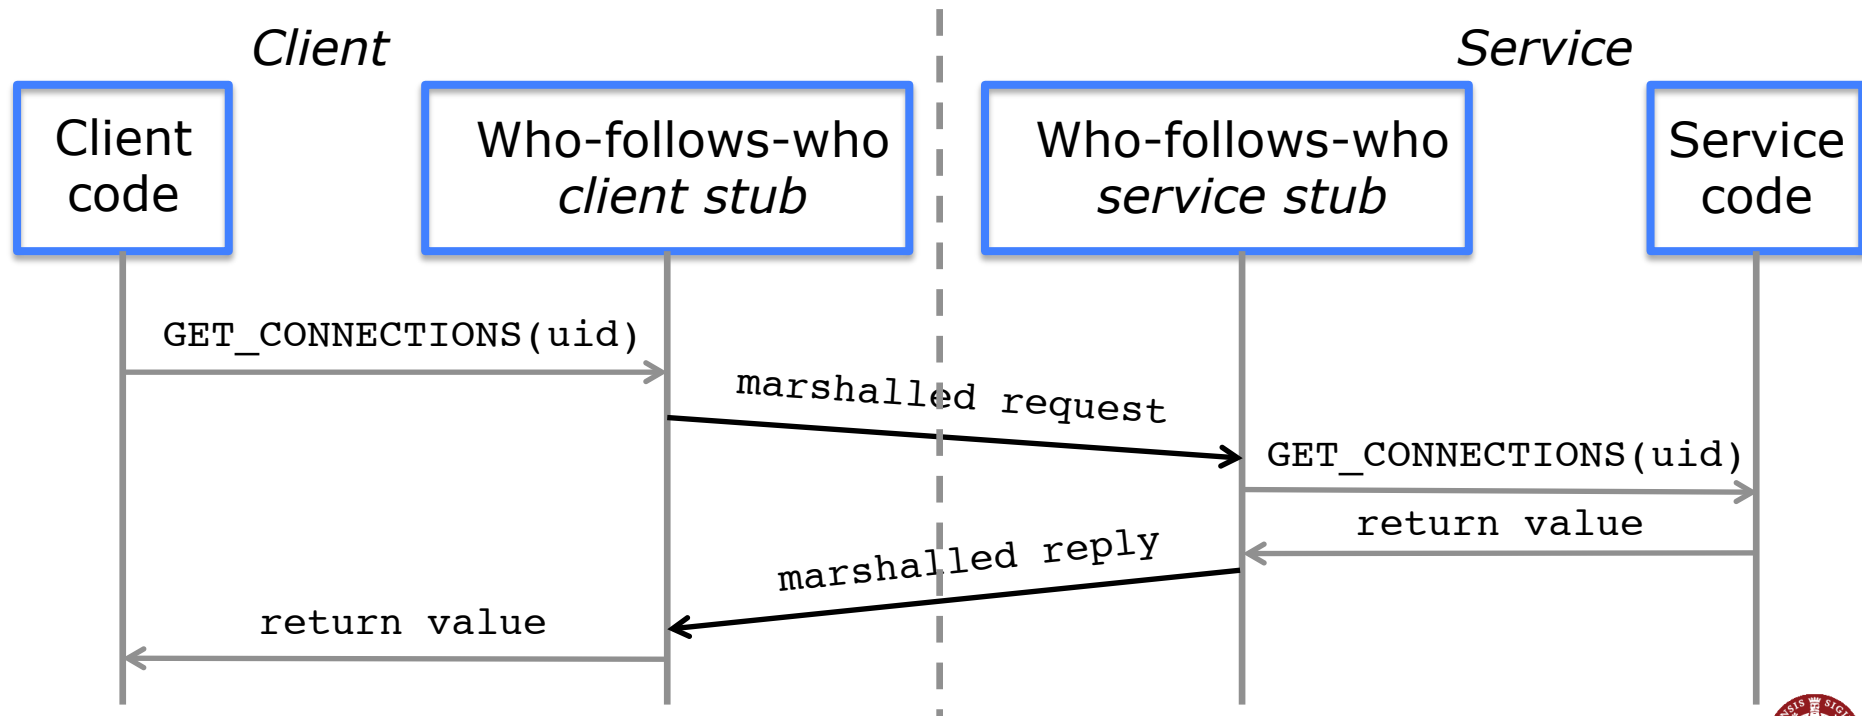
- Create virtualized versions of fundamental abstractions
- Client and services remain isolated even on same computer
- VMs: Virtualize the virtualizer ☺

Even more techniques for fault tolerance
later in the course



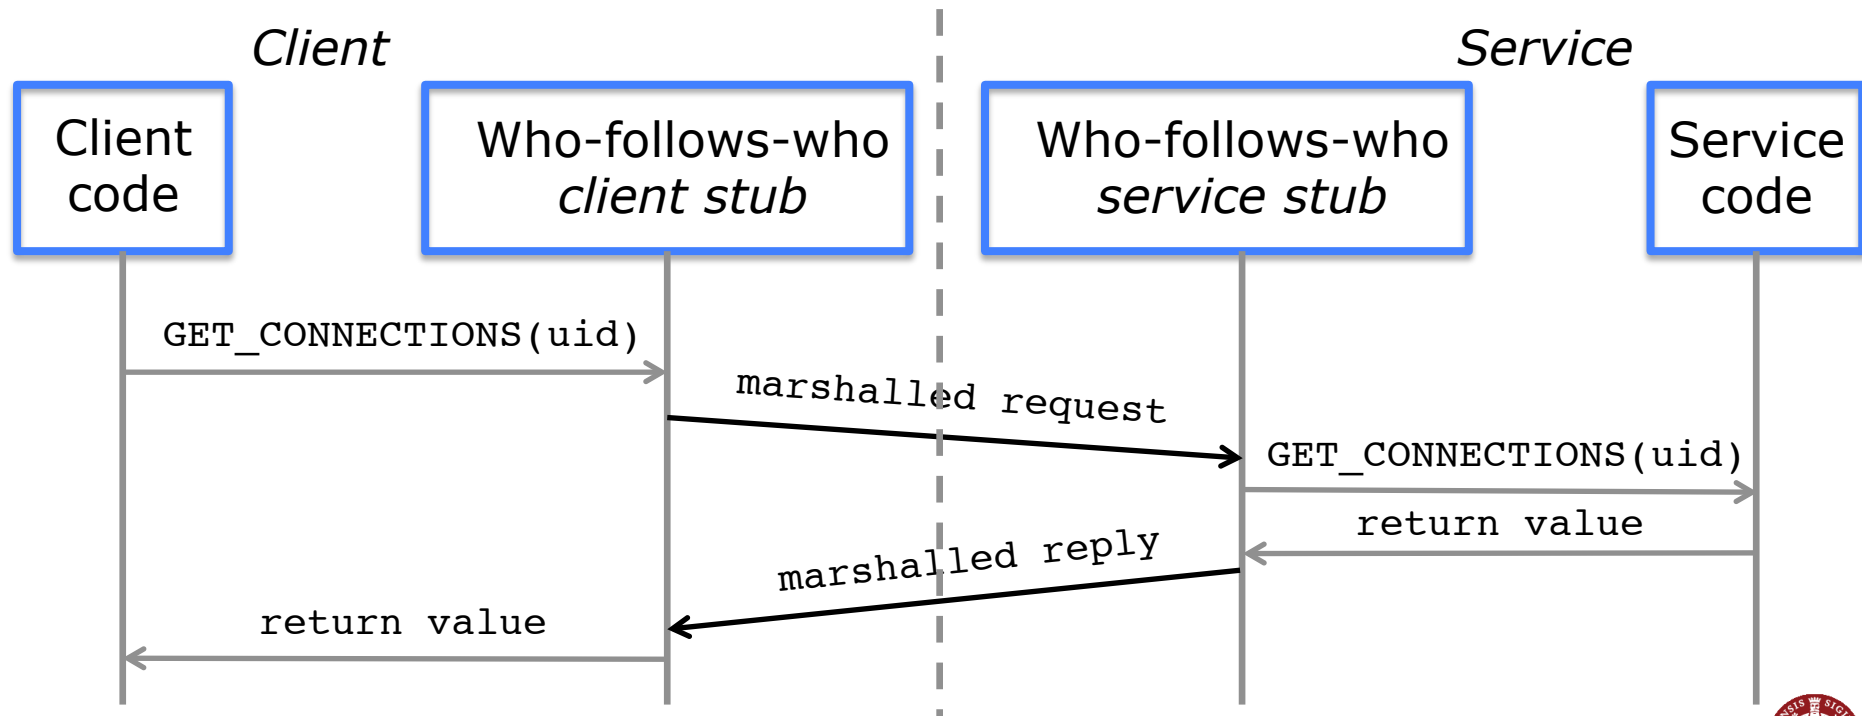
RPC: Remote Procedure Call

- Client-service request / response interactions
- Automate *marshalling* and *communication*



RPC: Remote Procedure Call

- How do RPC semantics differ from local procedure calls?
- What can go wrong? How can you fix it?



RPC Semantics

- **At-least-once**
 - Operation is *idempotent*
 - Naturally occurs if side-effect free
 - Stub just retries operation → failures can still occur!
 - Example: calculate SQRT
- **At-most-once**
 - Operation does have side-effects
 - Stub must ensure duplicate-free transmission
 - Example: transfer \$100 from my account to yours
- **Exactly-once**
 - Possible for certain classes of failures
 - Stub & service keep track (*durably*) of requests and responses
 - Example: bank cannot develop amnesia! 😊




How to achieve RPCs?

- Special-purpose ***request-reply protocol***, e.g., as in DNS
 - Developer must design protocol and marshalling scheme
- ***Classic RPC*** protocols, e.g., DCE, Sun RPC
 - Special APIs and schemes for marshalling
- ***RMI: Remote Method Invocation***
 - RPCs for methods in OO languages
 - Compiler-generated proxies
- ***Web Services***
 - Many modes of communication possible, including RPC-style communication
 - Tools available to compile proxies, e.g., JAX-WS
 - Generic marshalling (e.g., XML, JSON, Protocol Buffers) over HTTP transport → ***programming-language independence!***



RPC and Naming

- Most basic extension to the synchronous interaction pattern
 - Avoid having to name the destination
 - Ask where destination is
 - Then bind to destination
- Advantages:
 - Development is independent of deployment properties (e.g., network address)
 - More flexibility:
 - Change of address
 - Can be combined with:
 - Load balancing
 - Monitoring
 - Routing
 - Advanced service search



Source: Gustavo Alonso,
ETH Zurich (partial)



RPC Example: Domain Name Service (DNS)

- DNS: Translates domain names into IP addresses
 - `www.diku.dk` → `130.225.96.3` →
`d.root-servers.net` → `128.8.10.90` →
`p.nic.dk` → `204.61.216.36` →
`dns2.diku.dk` → `130.225.96.4` →
`web-aggregator.diku.dk` → `130.225.96.108`
- DNS typically implemented as a request-reply service, i.e., equivalent to RPC at high level



RPC Example: Domain Name Service (DNS)

- ***DNS interface***

RESOLVE(**domain_name**, **recursive**) → **value**

- ***DNS Implementation***

- DNS does not keep whole name database in a single server → **hierarchy of servers**
 - Not surprisingly, integrates **naming** into the service itself! 😊
- DNS uses **delegation**
 - Repeated RPCs to different servers
 - Recursive resolution and caching on some servers




DNS Operation and Caching

Source:
Kurose &
Ross
(partial)

- host at `cis.poly.edu` wants IP address for `gaia.cs.umass.edu`

Iterated query

- contacted server replies with name of server to contact
- **Caching** at local server can improve performance *significantly*
- But what about **semantics**?



Common Issues in Designing Services

- **Consistency**
 - How to deal with *updates* from multiple clients?
- **Coherence**
 - How to refresh *caches* while respecting consistency?
- **Scalability**
 - What happens to resource usage if we increase the #clients or the #operations?
- **Fault Tolerance**
 - Under what circumstances will the service be unavailable?



Other Examples of Services

- File systems: NFS, GFS
- Object stores: Dynamo, PNUTS
- Databases: pick your favorite relational DB ☺
- Configuration: Zookeeper
- Even whole computing clouds!
 - Infrastructure-as-a-service (IaaS), e.g., Amazon EC2, Rackspace, Windows Azure
 - Platform-as-a-service (PaaS), e.g., Windows Azure, Google AppEngine
 - Software-as-a-service (SaaS), e.g., Salesforce.com, Gmail
- And many, many others
- Differences in semantics are significant!



Questions so far?



Abstractions, Implementation and Performance

- Let I_1 and I_2 be two implementations of an abstraction
- **Examples**
 - Web service with or without HTTP proxies
 - Virtual memory with or without paging
 - Transactions via concurrency or serialization

How can we choose between I_1 and I_2 ?



Performance Metrics

latency
throughput
scalability
overhead
utilization
capacity

Discussion: What do these metrics mean?




Common Issues with Performance Metrics

- Properties of ***machines*** vs. properties of ***programs***
 - Utilization, capacity
 - Overhead, throughput
 - Latency, scalability
- Relationship between ***latency*** and ***throughput***
 - In serial case: $\text{latency} = 1 / \text{throughput}$
 - Not true when there is concurrency!



Performance and Hardware Trends

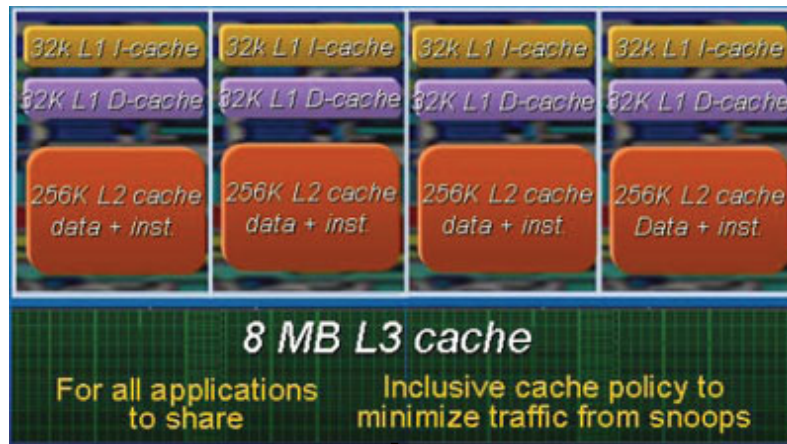
- OS-provided illusion of a computer:




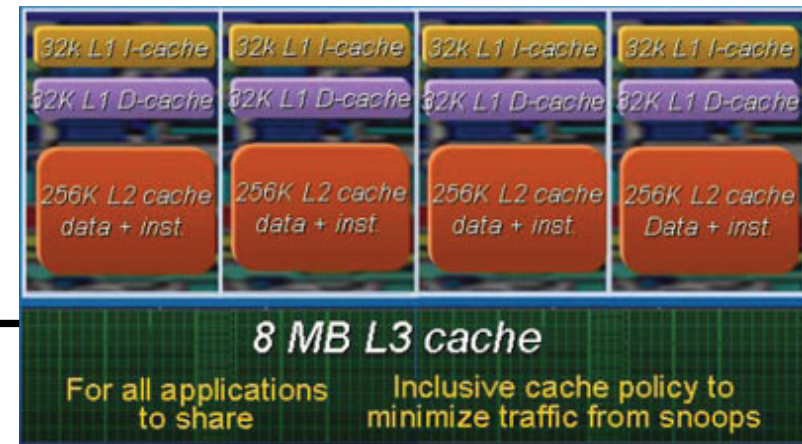
Does this picture look like a real computer?

How about this one?

Processor 1



Processor n



But the picture is not to scale!

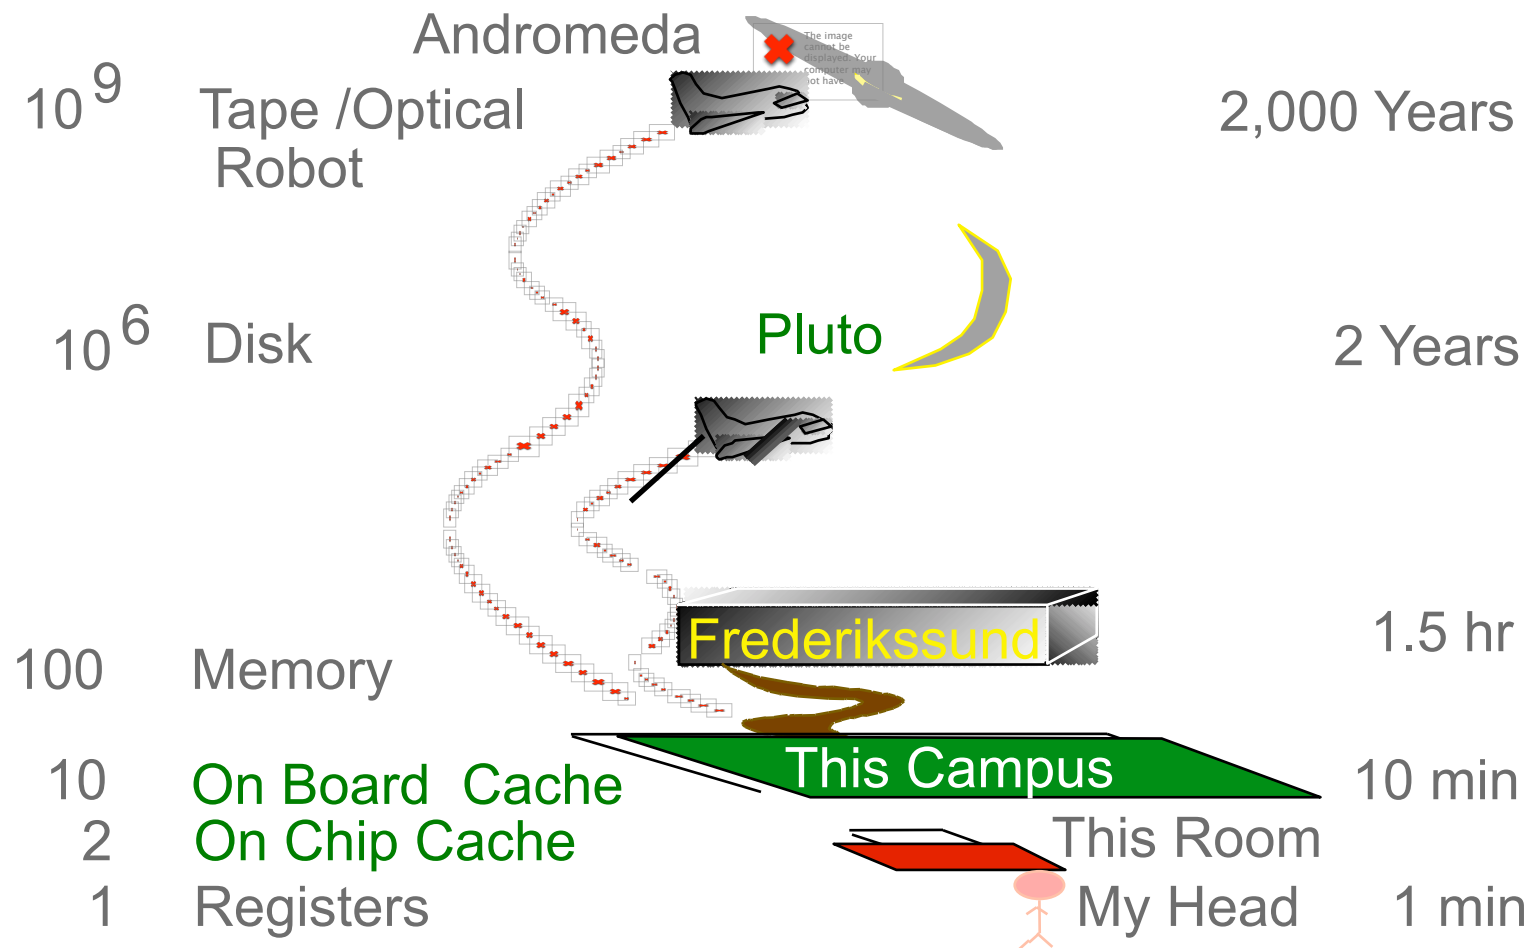
Size of
memory

Size of
last-level cache

What about the size of disk?



Storage Hierarchy




Source: Gray (partial)



And only getting worse...

- Riding Moore's Law
 - CPU clock speed
 - Memory size
 - Memory bandwidth
 - Disk size
 - Disk bandwidth
- Going way slower
 - Memory latency
 - Disk latency
- What does that do to random accesses?




Source: Boncz (partial)



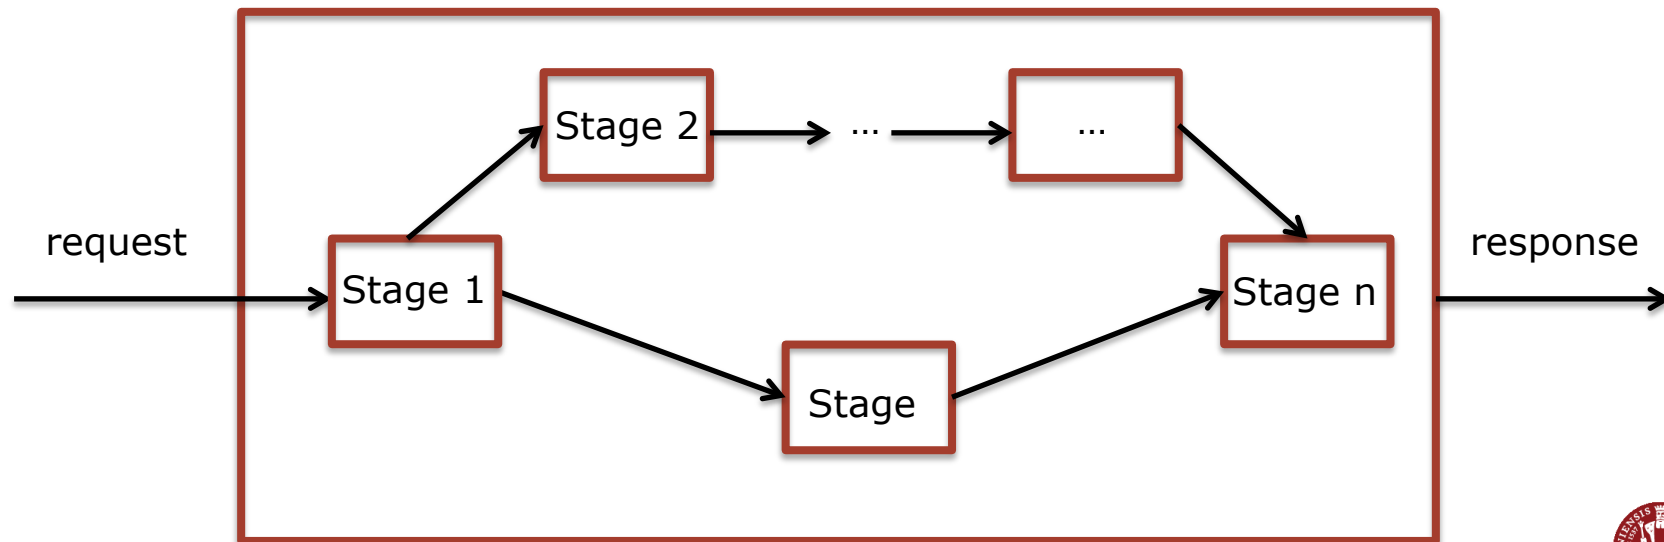
RAM = NQSRAM?!

- What we call
Random **A**ccess **M**emory
actually behaves as
Not-**Q**uite-**S**o-**R**andom **A**ccess **M**emory
because of the memory hierarchy
- Access to nearby cell **much faster** than to a far away cell!



How can we improve performance?

- Fast-path coding
 - Split processing into two code paths
 - One optimized path for common requests → fast path
 - One slow but comprehensive path for all other requests → slow path
 - Caching is an example of fast-path coding




Source: Saltzer & Kaashoek & Morris (partial)



How can we improve performance?


- Concurrency
 - Run multiple requests in different threads
 - Example: different web requests run in different threads or even servers
 - May improve both throughput **and** latency, but must be careful with locking, correctness
 - Can be hidden under abstractions, e.g., MapReduce and transactions



Remember RPC?

How can we improve performance?

- Batching
 - Run multiple requests at once
 - Example: batch I/Os and use elevator algorithm
 - May improve latency **and** throughput
- Dallying
 - Wait until you accumulate some requests and then run them
 - Example: group commit
 - May improve throughput when used together with batching, but typically incurs a latency penalty




How can we improve performance?

- Speculation, i.e., predict the future 😊
 - Guess the next requests and run them in advance
 - Example: prefetching
 - May overlap expensive operations, instead of waiting for their completion



Sounds good with reads, but can you speculate writes?

What should we learn today?



- Recognize and explain modular designs with clients and services
- Predict the functioning of service calls under different RPC semantics and failure modes
- Identify different mechanisms to achieve RPCs
- Implement RPC services with an appropriate mechanism, such as web services
- Explain performance metrics such as latency, throughput, overhead, utilization, capacity, and scalability
- List common hardware parameters that affect performance
- Apply performance improvement techniques, such as concurrency, batching, dallying, and fast-path coding

