



Concurrency Control: 2PL

Concurrency Control: Introduction to Schedules and Serializability

PCSD, Marcos Vaz Salles

Do-it-yourself Recap: Techniques for Performance

- What is the meaning of the following performance metrics: throughput, latency, overhead, utilization, capacity?
- Why can concurrency improve throughput and latency? How does that related to modern hardware characteristics?

latency
throughput
scalability
overhead
utilization
capacity



What should we learn today?



- Identify the multiple interpretations of the property of atomicity
- Implement methods to ensure before-or-after atomicity, and argue for their correctness
- Explain the variants of the two-phase locking (2PL) protocol, in particular the widely-used Strict 2PL
- Discuss definitions of serializability and their implications, in particular conflict-serializability and view-serializability
- Apply the conflict-serializability test using a precedence graph to transaction schedules



Read-Write Systems

- On-Line Transaction Processing (**OLTP**)
 - Process multiple, but relatively simple, application functions
- **Examples**
 - Order processing, e.g., Amazon
 - Item buy/sell in computer games, e.g., EVE Online
 - High-performance trading
 - Updates on social networks, e.g., Facebook



Atomicity vs. Performance is the fundamental trade-off

- Last week's property:
Strong Modularity
- This week:
Atomicity (before-or-after flavor)





Transaction

- Reliable unit of work against memory abstraction
 - In the next lectures, we will use “memory state” and “database” interchangeably!
- **ACID Properties**
 - **Atomicity**: transactions are all-or-nothing
 - **Consistency**: transaction takes database from one consistent state to another
 - **Isolation**: transaction executes as if it were the only one in the system (aka before-or-after atomicity)
 - **Durability**: once transaction is done (“committed”), results are persistent in the database



Examples of Transactions in SQL

Transaction T1: TRANSFER

```
BEGIN
  UPDATE account
  SET bal = bal + 100
  WHERE account_id = 'A';
  --
  UPDATE account
  SET bal = bal - 100
  WHERE account_id = 'B';
COMMIT
```

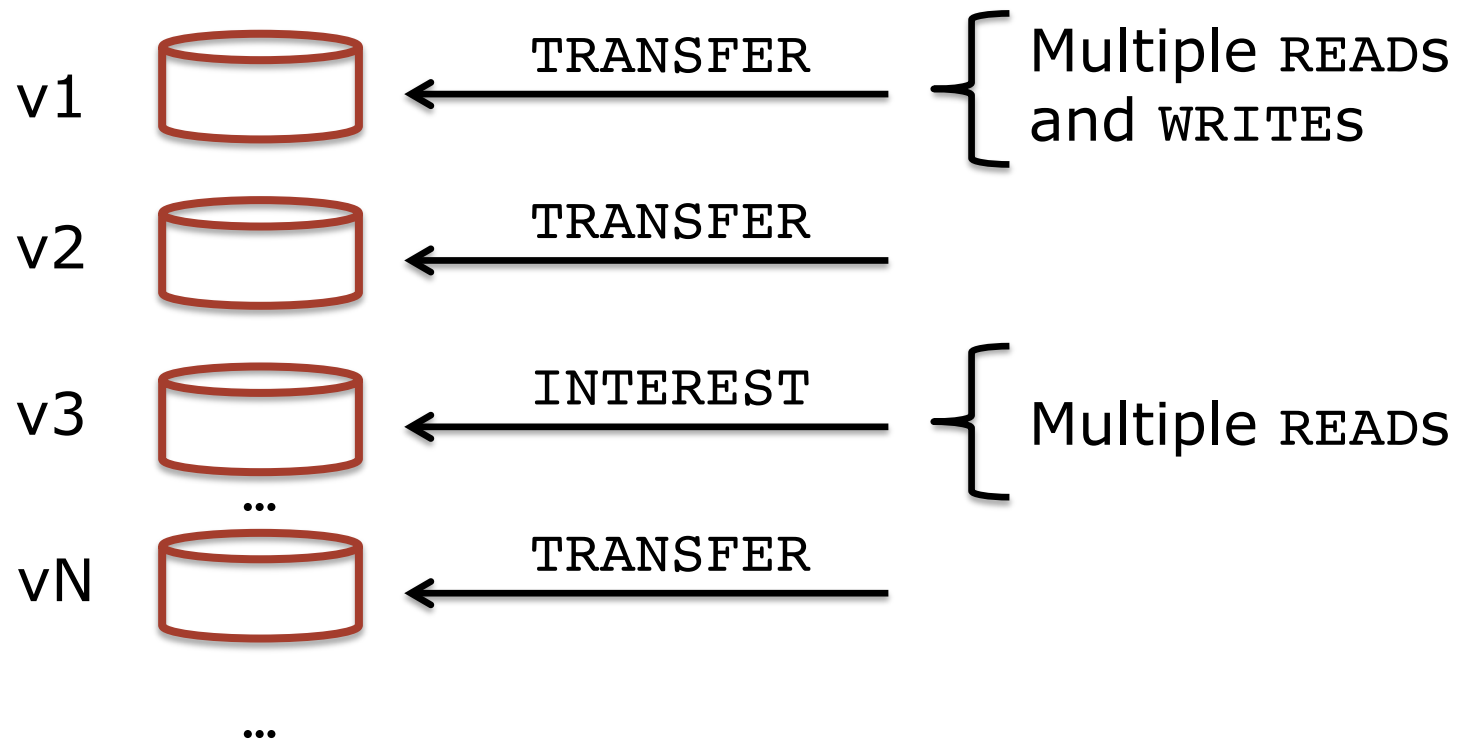
Transaction T2: INTEREST

```
BEGIN
  UPDATE account
  SET bal = bal * 1.06;
COMMIT
```

Under the hood, we
know it all translates to
calls to READ and WRITE



Conceptual Model: Version Histories




The many faces of atomicity

- **Atomicity** is strong modularity mechanism!
 - Hides that one high-level action is actually made of many sub-actions
- **Before-or-after** atomicity
 - == Isolation
 - Cannot have effects that would only arise by interleaving of parts of transactions
- **All-or-nothing** atomicity
 - == Atomicity (+ Durability)
 - Cannot have partially executed transactions
 - Once executed and confirmed, transaction effects are visible and not forgotten



Scaling Up



- **Problem:** Ensure automatically that all interactions leave data consistent



Goal of Concurrency Control

- Transactions should be executed so that it is *as though* they executed in some serial order
 - Also called **Isolation** or **Serializability** or **Before-or-after atomicity**
- Weaker variants also possible
 - Lower “degrees of isolation”



Example

- Consider again our two transactions (*Xacts*):

T1:	BEGIN	$A=A+100,$	$B=B-100$	END
T2:	BEGIN	$A=1.06*A,$	$B=1.06*B$	END

- T1** transfers \$100 from B's account to A's account
- T2** credits both accounts with 6% interest
- If submitted concurrently, net effect should be equivalent to *Xacts* running in some serial order
 - No guarantee that **T1** “logically” occurs before **T2** (or vice-versa) – but one of them is true



Solution 1


- 1) Get exclusive lock on entire database
 - 2) Execute transaction
 - 3) Release exclusive lock
- Transactions execute in *critical section*
 - Serializability guaranteed because execution is serial!
 - Problems?



Solution 2

- 1) Get exclusive locks on *accessed* data items
- 2) Execute transaction
- 3) Release exclusive locks


- Greater concurrency
- Problems?



Solution 3

- 1) Get exclusive locks on data items that are *modified*; get shared locks on data items that are only *read*
- 2) Execute transaction
- 3) Release all locks

- Greater concurrency
- Conservative Strict Two Phase Locking (2PL)
- Problems?



	S	X
S	Yes	No
X	No	No




Solution 4

- 1) Get exclusive locks on data items that are modified and get shared locks on data items that are read
- 2) Execute transaction and release locks on objects no longer needed *during execution*

- Greater concurrency
- Conservative Two Phase Locking (2PL)

- Problems?



	S	X
S	Yes	No
X	No	No




Solution 5

- 1) Get exclusive locks on data items that are modified and get shared locks on data items that are read, but do this *during execution* of transaction (as needed)
- 2) Release all locks

- Greater concurrency
- Strict Two Phase Locking (2PL)

- Problems?




	S	X
S	Yes	No
X	No	No



Solution 6

- 1) Get exclusive locks on data items that are modified and get shared locks on data items that are read, but do this during execution of transaction (as needed)
- 2) Release locks on objects no longer needed during execution of transaction
- 3) *Cannot acquire locks once any lock has been released***
 - ***Hence two-phase (acquiring phase and releasing phase)***
- Greater concurrency
- Two Phase Locking (2PL)
- Problems?



Summary of Alternatives


- Conservative Strict 2PL
 - No deadlocks, no cascading aborts
 - **But** need to know objects a priori, least concurrency
- Conservative 2PL
 - No deadlocks, more concurrency than Conservative Strict 2PL
 - **But** need to know objects a priori, when to release locks, cascading aborts
- Strict 2PL
 - No cascading aborts, no need to know objects a priori or when to release locks, more concurrency than Conservative Strict 2PL
 - **But** deadlocks
- 2PL
 - Most concurrency, no need to know object a priori
 - **But** need to know when to release locks, cascading aborts, deadlocks

Source: Ramakrishnan & Gehrke (partial)



Method of Choice

- Strict 2PL
 - No cascading aborts, no need to know objects a priori or when to release locks, more concurrency than Conservative Strict 2PL
 - But deadlocks
- Reason for choice
 - Cannot know objects a priori, so no Conservative options
 - Thus only 2PL and Strict 2PL left
 - 2PL needs to know when to release locks (main problem), and has cascading aborts
 - Hence Strict 2PL
- Implication
 - Need to deal with deadlocks!



	S	X
S	Yes	No
X	No	No



Lock Management

- Lock/unlock requests handled by lock manager
- Lock table entry:
 - Number of transactions currently holding a lock
 - Type of lock held (shared or exclusive)
 - Pointer to queue of lock requests
- Locking and unlocking have to be atomic operations
- **Lock upgrade:** transaction that holds a shared lock can be upgraded to hold an exclusive lock



Questions so far?



Is Strict 2PL correct?

- We will formalize now **serializability** and argue that Strict 2PL is correct
 - Full proof is left as homework 😊
- Strict 2PL can however deadlock
 - We will see how to handle deadlock automatically



Schedules

- Consider a possible interleaving (schedule):

T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.06 * A, \quad B = 1.06 * B$	

- The system's view of the schedule:

T1:	$R(A), W(A),$	$R(B), W(B)$
T2:	$R(A), W(A), R(B), W(B)$	



Scheduling Transactions

- *Serial schedule*: Schedule that does not interleave the actions of different transactions.
- *Equivalent schedules*: For any database state
 - The effect (on the set of objects in the database) of executing the schedules is the same
 - The values read by transactions is the same in the schedules
 - Assume no knowledge of transaction logic
- *Serializable schedule*: A schedule that is equivalent to some serial execution of the transactions.

(Note: If each transaction preserves consistency, every serializable schedule preserves consistency.)



Anomalies with Interleaved Execution

- Reading Uncommitted Data (WR Conflicts, “dirty reads”):

T1:	R(A), W(A),	R(B), W(B), Abort
T2:	R(A), W(A), C	

- Unrepeatable Reads (RW Conflicts):

T1:	R(A),	R(A), W(A), C
T2:	R(A), W(A), C	



Anomalies (contd.)

- Overwriting Uncommitted Data (WW Conflicts):

T1:	W(A),	W(B), C
T2:	W(A), W(B), C	



Conflict Serializable Schedules


- Two schedules are **conflict equivalent** if:
 - Involve the same actions of the same transactions
 - Every pair of conflicting actions is ordered the same way
- Schedule S is **conflict serializable** if S is conflict equivalent to some serial schedule



Example

- A schedule that is not conflict serializable:

T1:	R(A), W(A),	R(B), W(B)
T2:	R(A), W(A), R(B), W(B)	



- The cycle in the graph reveals the problem. The output of T1 depends on T2, and vice-versa.



Precedence Graph

- Precedence graph: One node per T_i ; edge from T_i to T_j if operation in T_j conflicts with earlier operation in T_i .
- Theorem: Schedule is conflict serializable if and only if its precedence graph is acyclic
- Strict 2PL only allows conflict serializable schedule
 - Precedence graph is always acyclic



Are the following schedules conflict-serializable?

- Build the precedence graph for each of the following transaction schedules

T1:	R(A)		W(B)			C
T2:		R(B)		R(A)	R(C)	C
T3:			R(B)		W(C)	C

Note: C stands for commit

T1:	R(A)		W(B)			C
T2:		R(B)		R(A)	R(C)	C
T3:			R(B)		W(C)	C



Returning to Definition of Serializability

- A schedule S is serializable if there exists a serial order SO such that:
 - The **state of the database** after S is the **same** as the state of the database after SO
 - The **values read** by each transaction in S are the **same** as that returned by each transaction in SO
 - Database does not know anything about the internal structure of the transaction programs
- **Under this definition, certain serializable executions are not conflict serializable!**



Are these schedules serializable?

T1:	R(A)		W(A)
T2:		W(A)	
T3:			W(A)

T1:	R(A), W(A)		
T2:		W(A)	
T3:			W(A)



View Serializability

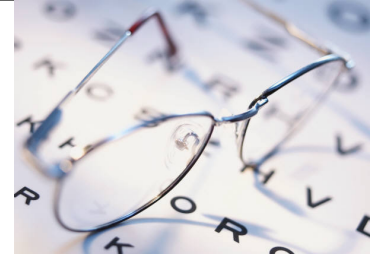
- Schedules S1 and S2 are **view equivalent** if:
 - If T_i reads initial value of A in S1, then T_i also reads initial value of A in S2
 - If T_i reads value of A written by T_j in S1, then T_i also reads value of A written by T_j in S2
 - If T_i writes final value of A in S1, then T_i also writes final value of A in S2

T1:	R(A)		W(A)
T2:		W(A)	
T3:			W(A)

T1:	R(A), W(A)		
T2:		W(A)	
T3:			W(A)



What should we learn today?



- Identify the multiple interpretations of the property of atomicity
- Implement methods to ensure before-or-after atomicity, and argue for their correctness
- Explain the variants of the two-phase locking (2PL) protocol, in particular the widely-used Strict 2PL
- Discuss definitions of serializability and their implications, in particular conflict-serializability and view-serializability
- Apply the conflict-serializability test using a precedence graph to transaction schedules

