



# Concurrency Control: Serializability, Schedules, Advanced Topics

PCSD, Marcos Vaz Salles

# Do-it-yourself-recap: Locking Solutions for Isolation in ACID Transactions

	S	X
S	Yes	No
X	No	No

Conflict or View Serializable? Deadlocks? Cascading aborts?


## Solution 4

- 1) Get exclusive locks on data items that are modified and get shared locks on data items that are read
- 2) Execute transaction and release locks on objects no longer needed *during execution*

- Greater concurrency
- Conservative Two Phase Locking (2PL)

- Problems?

2




## Solution 5

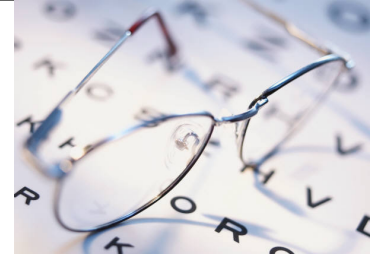
- 1) Get exclusive locks on data items that are modified and get shared locks on data items that are read, but do this *during execution* of transaction (as needed)
- 2) Release all locks

- Greater concurrency
- Strict Two Phase Locking (2PL)

- Problems?



## What should we learn today?



- Explain deadlock prevention and detection techniques
- Apply deadlock detection using a waits-for graph to transaction schedules
- Explain situations where predicate locking is required
- Explain the optimistic concurrency control and multi-version concurrency control models
- Predict validation decisions under optimistic concurrency control



## Deadlocks

- Deadlock: Cycle of transactions waiting for locks to be released by each other.
- Two ways of dealing with deadlocks:
  - Deadlock prevention
  - Deadlock detection



## Deadlock Prevention

- Assign priorities based on timestamps. Assume  $T_i$  wants a lock that  $T_j$  holds. Two policies are possible:
  - Wait-Die: If  $T_i$  has higher priority,  $T_i$  waits for  $T_j$ ; otherwise  $T_i$  aborts
  - Wound-wait: If  $T_i$  has higher priority,  $T_j$  aborts; otherwise  $T_i$  waits
- If a transaction re-starts, make sure it has its original timestamp



## Deadlock Detection



- Create a **waits-for graph**:
  - Nodes are transactions
  - There is an edge from  $T_i$  to  $T_j$  if  $T_i$  is waiting for  $T_j$  to release a lock
- Periodically check for cycles in the waits-for graph



# Deadlock Detection

## • Example

T1: S(A), R(A), S(B)  
 T2: X(B), W(B) X(C)  
 T3: S(C), R(C) X(A)  
 T4: X(B)



Do the following schedules lead to deadlock?

- Build the waits-for graph for each of the following transaction schedules

T1:	S(A)		X(D)	X(C)	C	
T2:			X(A)			X(B)
T3:		S(B)				S(C)

Note: we only show locking operations for brevity! C denotes commit, all locks released

T1:			S(C)	S(A)		X(D)
T2:	S(B)				X(C)	
T3:		S(D)	X(B)			





## Dynamic Databases: Locking the objects that exist now in the database is not enough!

- If we relax the assumption that the DB is a fixed collection of objects, even Strict 2PL will not assure serializability:
  - T1 locks all pages containing sailor records with *rating* = 1, and finds oldest sailor (say, *age* = 71).
  - Next, T2 inserts a new sailor; *rating* = 1, *age* = 96.
  - T2 also deletes oldest sailor with *rating* = 2 (and, say, *age* = 80), and commits.
  - T1 now locks all pages containing sailor records with *rating* = 2, and finds oldest (say, *age* = 63).
- No consistent DB state where T1 is “correct”!




## The Problem

- T1 implicitly assumes that it has locked the set of all sailor records with *rating* = 1.
  - Assumption only holds if no sailor records are added while T1 is executing!
  - Need some mechanism to enforce this assumption.  
(Index locking and predicate locking.)
- Example shows that conflict serializability guarantees serializability only if the set of objects is fixed!




## Index Locking

- If data is accessed by an **index** on the *rating* field, T1 should **lock the index page** containing the data entries with *rating* = 1.
  - If there are no records with *rating* = 1, T1 must lock the index page where such a data entry *would* be, if it existed!
- If there is **no suitable index**, T1 must **lock all pages**, and lock the file/table to prevent new pages from being added, to ensure that no new records with *rating* = 1 are added.



## Multiple-Granularity Locks

- Hard to decide what granularity to lock (tuples vs. pages vs. tables).
- Shouldn't have to decide!
- Data “containers” are nested:




## Solution: New Lock Modes, Protocol

- Allow Xacts to lock at each level, but with a special protocol using new **"intention" locks**
- Before locking an item, Xact must set "intention locks" on all its ancestors.
- For unlock, go from specific to general (i.e., bottom-up).
- **SIX mode:** Like S & IX at the same time.

	--	IS	IX	S	X
--	✓	✓	✓	✓	✓
IS	✓	✓	✓	✓	
IX	✓	✓	✓		
S	✓	✓		✓	
X	✓				



## Examples: Multiple-Granularity Locks



### Scenario 1:


T1 scans Emp;  
T2 uses  
indirect index,  
finds e3

**T1:** SELECT \* FROM Emp  
WHERE age > 25  
**T2:** UPDATE Emp  
SET sal = sal \* 1.1  
WHERE ssn = 42

	--	IS	IX	S	X
--	✓	✓	✓	✓	✓
IS	✓	✓	✓	✓	
IX	✓	✓	✓		
S	✓	✓		✓	
X	✓				



## Examples: Multiple-Granularity Locks



### Scenario 2:

T1 and T2 use indexes; T1 starts range scan on e1, T2 finds e3

**T1:** SELECT \* FROM Emp  
WHERE age > 25

**T2:** UPDATE Emp  
SET sal = sal \* 1.1  
WHERE ssn = 42

T1:ST1:ST2:X  
~~T1:S~~  
T1  
waits

	--	IS	IX	S	X
--	✓	✓	✓	✓	✓
IS	✓	✓	✓	✓	
IX	✓	✓	✓		
S	✓	✓		✓	
X	✓				



Questions so far?





## The Problems with Locking


- Locking is a pessimistic approach in which conflicts are prevented. Disadvantages:
  - Lock management overhead.
  - Deadlock detection/resolution.
  - Lock contention for heavily used objects.
- Remember: We must devise a way to **enforce serializability**, without destroying concurrency
- Two approaches:
  - Prevent violations → **locking**
  - Fix violations → **aborts**

How can we design a protocol based on  
aborts instead of locks?



## Optimistic CC: Kung-Robinson Model

- Xacts have three phases
- **READ**: Xacts read from the database, but make changes to private copies of objects.
- **VALIDATE**: Check for conflicts.
- **WRITE**: Make local copies of changes public.



## Validation

- Test conditions that are **sufficient** to ensure that no conflict occurred.
- Each Xact is assigned a numeric id.
  - Just use a **timestamp**.
- Xact ids assigned at end of READ phase, just before validation begins. (Why then?)
- **ReadSet(Ti)**: Set of objects read by Xact Ti.
- **WriteSet(Ti)**: Set of objects modified by Ti.



## Test 1

- For all  $i$  and  $j$  such that  $T_i < T_j$ , check that  $T_i$  completes before  $T_j$  begins.



## Test 2

- For all  $i$  and  $j$  such that  $T_i < T_j$ , check that:
  - $T_i$  completes before  $T_j$  begins its Write phase +
  - $\text{WriteSet}(T_i) \cap \text{ReadSet}(T_j)$  is empty.




Does  $T_j$  read dirty data? Does  $T_i$  overwrite  $T_j$ 's writes?



## Test 3

- For all  $i$  and  $j$  such that  $T_i < T_j$ , check that:
  - $T_i$  completes Read phase before  $T_j$  does +
  - $\text{WriteSet}(T_i) \cap \text{ReadSet}(T_j)$  is empty +
  - $\text{WriteSet}(T_i) \cap \text{WriteSet}(T_j)$  is empty.



Does  $T_j$  read dirty data? Does  $T_i$  overwrite  $T_j$ 's writes?



## Validation Example

- Predict whether T3 will be allowed to commit, given the transactions below

T1:  $RS(T1) = \{1, 2, 3\}$ ,  $WS(T1) = \{4\}$ ,  
T1 completes before T3 begins with its write phase.  
T2:  $RS(T2) = \{6, 7, 8\}$ ,  $WS(T2) = \{8\}$ ,  
T2 completes read phase before T3 does.  
T3:  $RS(T3) = \{3, 5, 6, 7\}$ ,  $WS(T3) = \{4\}$ ,  
allow commit or roll back?



## Overheads in Optimistic CC

- Must record read/write activity in ReadSet and WriteSet per Xact.
  - Must create and destroy these sets as needed.
- Must check for conflicts during validation, and must make validated writes “global”.
  - Critical section can reduce concurrency.
  - Scheme for making writes global can reduce clustering of objects.
- Optimistic CC restarts Xacts that fail validation.
  - Work done so far is wasted; requires clean-up.
- Still, optimistic techniques widely used in software transactional memory (STM)





## Multiversion Concurrency Control (MVCC)

- This approach maintains a number of **versions** of a data item and allocates the **right version to a read operation** of a transaction. Thus unlike other mechanisms a **read operation in this mechanism is never rejected**.
- Side effect:
  - Significantly more storage (RAM and disk) is required to maintain multiple versions. To check unlimited growth of versions, a **garbage collection** is run when some criteria is satisfied
- Many commercial database systems implement a combination of MVCC and S2PL
- See compendium for more details



## Snapshot Isolation

- Often databases implement properties that are **weaker** than serializability
- **Snapshot isolation**
  - **Snapshots:** Transactions see snapshot as of beginning of their execution
  - **First Committer Wins:** Conflicting writes to same item lead to aborts
- May lead to **write skew**
  - Database must have at least one doctor on call
  - Two doctors on call concurrently examine snapshot and see exactly each other on call
  - Doctors update their own records to being on leave
    - No write-write conflicts: different records!
  - After commits, database has no doctors on call



## Transaction Support in SQL-92

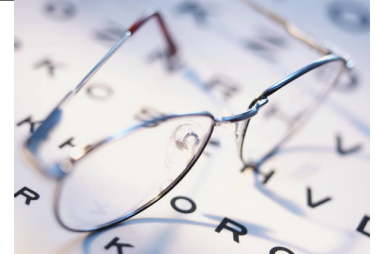
- Each transaction has an access mode, a diagnostics size, and an isolation level.

Does not correspond to serializability!

<b>Isolation Level</b>	<b>Dirty Read</b>	<b>Unrepeatable Read</b>	<b>Phantom Problem</b>
Read Uncommitted	Maybe	Maybe	Maybe
Read Committed	No	Maybe	Maybe
Repeatable Reads	No	No	Maybe
<b>Serializable</b>	No	No	No



## What should we learn today?



- Explain deadlock prevention and detection techniques
- Apply deadlock detection using a waits-for graph to transaction schedules
- Explain situations where predicate locking is required
- Explain the optimistic concurrency control and multi-version concurrency control models
- Predict validation decisions under optimistic concurrency control

