



# Recovery: ARIES, normal operation and crash recovery procedure

PCSD, Marcos Vaz Salles

## Do-it-yourself-recap: How to handle the Buffer Pool?

Recall:

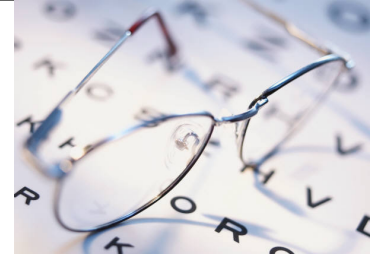
When do you need to UNDO changes?  
When do you need to REDO changes?

- **Force** every write to disk?
  - Poor response time.
  - But provides durability.
- **Steal** buffer-pool frames from uncommitted Xacts?
  - If not, poor throughput.
  - If so, how can we ensure atomicity?

	No Steal	Steal
Force	Trivial (?)	
No Force		Desired



## What should we learn today?



- Predict what portions of the log and database are necessary for recovery under different failure scenarios
- Explain how write-ahead logging is achieved in the ARIES protocol
- Explain the functions of recovery metadata such as the transaction table and the dirty page table
- Predict how recovery metadata is updated during normal operation
- Interpret the contents of the log resulting from ARIES normal operation
- Explain the three phases of ARIES crash recovery: analysis, redo, and undo
- Predict how recovery metadata, system state, and the log are updated during recovery




## Undo/Redo vs. Force/Steal

	No Steal	Steal
Force	No Redo No Undo	No Redo Undo
No Force	Redo No Undo	Redo Undo

How do we support this option?



## Basic Idea: Logging



- Record REDO and UNDO information, for every update, in a *log*.
  - Sequential writes to log (put it on a separate disk).
  - Minimal info (diff) written to log, so multiple updates fit in a single log page.
- Log: An ordered list of REDO/UNDO actions
  - Logical vs. Physical Logging
  - Example physical log record contains:

<XID, pageID, offset, length, old data, new data>

- Good compromise is physiological logging.





## Write-Ahead Logging (WAL)

- Golden Rule: Never modify the only copy!
- The **Write-Ahead Logging** Protocol:
  - 1) Must **force** the **log record** for an update before the corresponding **data page** gets to disk.
  - 2) Must **write all log records** for a Xact before commit.
- #1 guarantees Atomicity.
- #2 guarantees Durability.
- Exactly how is logging (and recovery!) done?
  - We will study the ARIES algorithms.




## Recovery Equations



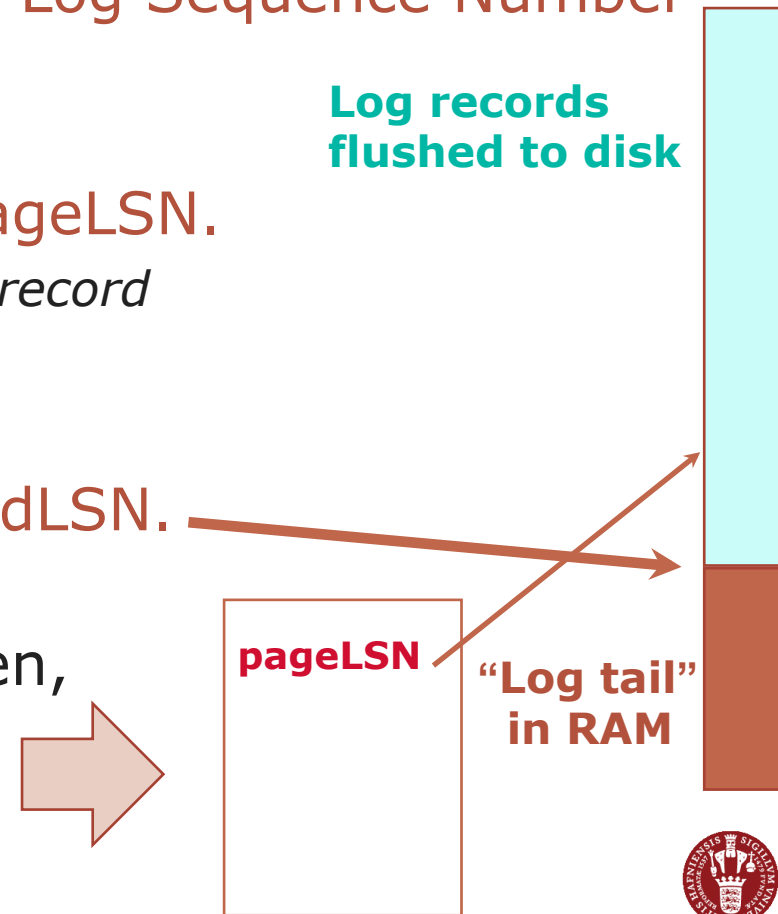
- Crash Recovery: volatile memory lost **Discussion:**
  - Current DB = DB files + DB log  $\longrightarrow$  since when?
- Media Recovery: nonvolatile storage lost
  - Current DB = DB backup + DB log  $\longrightarrow$  since when?
- We will focus on crash recovery next



## WAL & the Log




- Each log record has a unique **Log Sequence Number (LSN)**.
  - LSNs always increasing.
- Each data page contains a **pageLSN**.
  - The LSN of the most recent *log record* for an update to that page.
- System keeps track of **flushedLSN**.
  - The max LSN flushed so far.
- WAL: *Before* a page is written,
  - $\text{pageLSN} \leq \text{flushedLSN}$





## Log Records

### LogRecord fields:



Possible log record types:

- **Update**
- **Commit**
- **Abort**
- **End** (signifies end of commit or abort)
- **Compensation Log Records (CLRs)**
  - for UNDO actions

Note: Format above simplified; in reality, ARIES uses physiological variant

Source: Ramakrishnan & Gehrke (partial)



## Other Log-Related State

- **Transaction Table:**
  - One entry per active Xact.
  - Contains **XID**, **status** (running/committed/aborted), and **lastLSN**.
- **Dirty Page Table:**
  - One entry per dirty page in buffer pool.
  - Contains **recLSN** -- the LSN of the log record which ***first*** caused the page to be dirty.



## Normal Execution of an Xact

- Series of **reads** & **writes**, followed by **commit** or **abort**.
  - We will assume that write is atomic on disk.
    - In practice, additional details to deal with non-atomic writes.
- **Strict 2PL** → concurrency is correctly handled
- **STEAL, NO-FORCE** buffer management, with **Write-Ahead Logging**.



## Checkpointing

- Periodically, the DBMS creates a checkpoint, in order to minimize the time taken to recover in the event of a system crash. Write to log:
  - **begin\_checkpoint** record: Indicates when chkpt began.
  - **end\_checkpoint** record: Contains current *Xact table* and *dirty page table*. This is a `fuzzy checkpoint':
    - Other Xacts continue to run; so these tables accurate only as of the time of the **begin\_checkpoint** record.
    - No attempt to force dirty pages to disk; effectiveness of checkpoint limited by oldest unwritten change to a dirty page, **minDirtyPagesLSN**.
    - Use background process to flush dirty pages to disk!
  - Store LSN of chkpt record in a safe place (**master** record).




## The Big Picture: What's Stored Where



### LogRecords

prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image



### Data pages

each  
with a  
pageLSN

**master record**



### Xact Table

lastLSN  
status

### Dirty Page Table

recLSN

**flushedLSN**



## Transaction Commit

- Write **commit** record to log.
- All log records up to Xact's **lastLSN** are flushed.
  - Guarantees that **flushedLSN**  $\geq$  **lastLSN**.
  - Note that log flushes are sequential, synchronous writes to disk.
  - Many log records per log page.
- Commit() returns.
- Write **end** record to log.

Why?




## Simple Transaction Abort

- For now, consider an explicit abort of a Xact.
  - No crash involved.
- We want to “play back” the log in reverse order, UNDOing updates.
  - Get **lastLSN** of Xact from Xact table.
  - Can follow chain of log records backward via the **prevLSN** field.
  - Before starting UNDO, write an **Abort log record**.
    - For recovering from crash during UNDO!



## Abort, cont.



- To perform UNDO, must have a lock on data!
  - Strict 2PL enforces this
- Before restoring old value of a page, write a CLR:
  - You continue logging while you UNDO!!
  - CLR has one extra field: **undonextLSN**
    - Points to the next LSN to undo (i.e. the prevLSN of the record we're currently undoing).
  - CLRs *never* Undone (but they might be Redone when repeating history: guarantees Atomicity!)
- At end of UNDO, write an **end** log record.





## Example

- 10 T1 writes P5
- 20 T2 writes P17
- 30 T1 writes P3
- 31 P3 written to disk  
(pageLSN for page 3 at this time is 30)
- 40 T1 aborts
- 50 CLR T1 P3 (undonextLSN: 10)
- 60 CLR T1 P5 (undonextLSN: NULL)
- End T1



## A Longer Example

- 10 T1 writes P3 (prevLSN: NULL)
- 20 T2 writes P4 (prevLSN: NULL)
- 30 T2 writes P5 (prevLSN: 20)
- flushedLSN = 20

P4 gets written to disk (pageLSN for page 4 = 20)

T2 aborts

- 50 Abort T2
- 60 CLR T2 P5 (undoNextLSN = 20), pageLSN(P5)=60

Update P5 in the buffer manager

Flush log up to log record 60

Buffer manager writes P5 to disk.

- 70 CLR T2 P4 (undoNextLSN = NULL), pageLSN(P4)=70

Update page P4

- 80 End T2
- 90 T1 commits

Flush log up to log record 90, then the commit(T1) returns

Discussion: Does this example make sense?  
Can you explain to your colleague what happened?

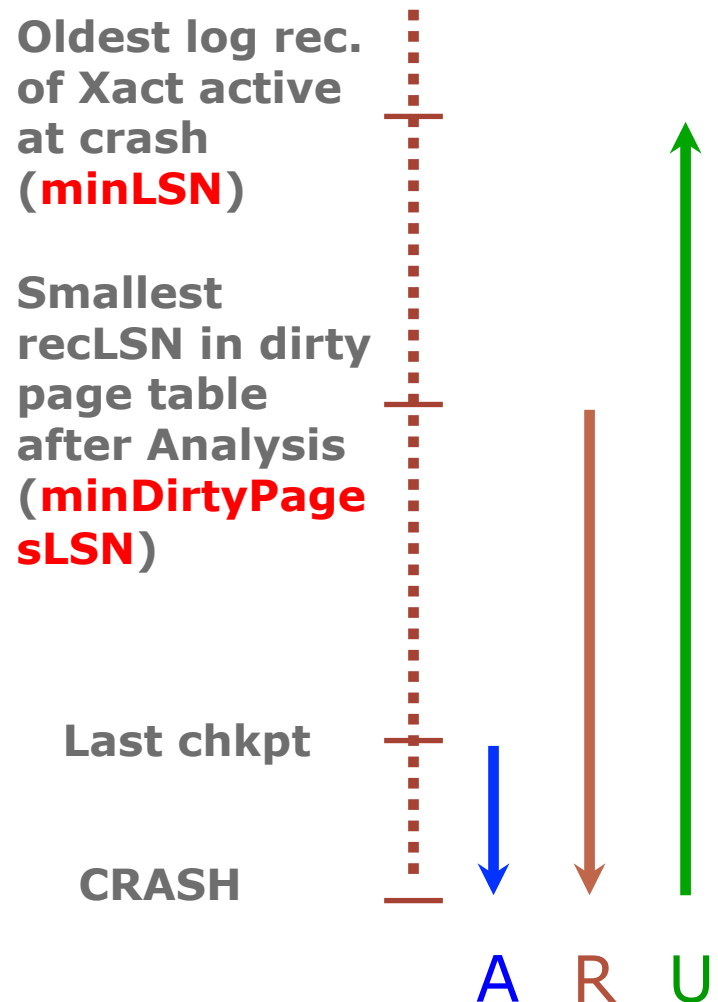


Questions so far?



# Crash Recovery: Big Picture

Keep in Mind:  
It must be OK to  
crash at **any time**  
during recovery



- ❖ Start from a **checkpoint** (found via **master** record).
- ❖ Three phases. Need to:
  - Figure out which Xacts committed since checkpoint, which failed (**Analysis**).
  - **REDO** *all* actions.
    - ♦ **Repeat History**
  - **UNDO** effects of failed Xacts.

Source: Ramakrishnan & Gehrke (partial)



## Recovery: The Analysis Phase

- Reconstruct state at checkpoint.
  - via **end\_checkpoint** record.
- Scan log forward from checkpoint.
  - **End** record: Remove Xact from Xact table.
  - **Other records**: Add Xact to Xact table, set **lastLSN=LSN**, change Xact status on **commit**.
  - **Update** record: If P not in Dirty Page Table,
    - Add P to D.P.T., set its **recLSN=LSN**.



## Recovery: The REDO Phase

- We *repeat History* to reconstruct state at crash:
  - Reapply *all* updates (even of aborted Xacts!), redo CLR.
- Scan forward from log rec containing smallest **recLSN** in D.P.T. For each CLR or update log rec **LSN**, REDO the action unless:
  - Affected page is not in the Dirty Page Table, or
  - Affected page is in D.P.T., but has **recLSN** > **LSN**, or
  - **pageLSN** (in DB) >= **LSN**.
- To **REDO** an action:
  - Reapply logged action.
  - Set **pageLSN** to **LSN**. No additional logging! (**Why?**)



## Recovery: The UNDO Phase

$ToUndo = \{ lsn \mid lsn \text{ a lastLSN of a "loser" Xact} \}$

### Repeat:

- Choose largest LSN among ToUndo.
- If this LSN is a CLR and  $undonextLSN == NULL$ 
  - Write an End record for this Xact.
- If this LSN is a CLR, and  $undonextLSN \neq NULL$ 
  - Add  $undonextLSN$  to ToUndo
- Else this LSN is an update. Undo the update, write a CLR, add  $prevLSN$  to ToUndo.

**Until ToUndo is empty.**



## Example of Recovery



Xact Table

lastLSN  
status

Dirty Page Table

recLSN

flushedLSN

ToUndo

LSN	LOG
00	begin_checkpoint
05	end_checkpoint
10	update: T1 writes P5
20	update T2 writes P3
30	T1 abort
40	CLR: Undo T1 LSN 10
45	T1 End
50	update: T3 writes P1
60	update: T2 writes P5
	CRASH, RESTART

prevLSNs





## Example: Crash During Restart!

Discussion:  
Explain to your  
colleague!



Xact Table

lastLSN

status

Dirty Page Table

recLSN

flushedLSN

ToUndo

LSN	LOG
00,05	begin_checkpoint, end_checkpoint
10	update: T1 writes P5
20	update T2 writes P3
30	T1 abort
40,45	CLR: Undo T1 LSN 10, T1 End
50	update: T3 writes P1
60	update: T2 writes P5
<del>70</del>	<del>CRASH, RESTART</del>
70	CLR: Undo T2 LSN 60
80,85	CLR: Undo T3 LSN 50, T3 end
<del>90</del>	<del>CRASH, RESTART</del>
90	CLR: Undo T2 LSN 20, T2 end

undonextLSN

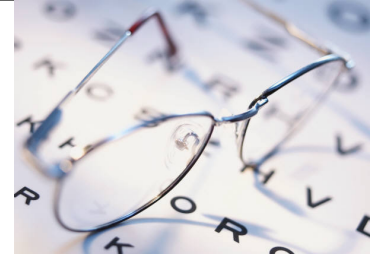


## Additional Crash Issues

- What happens if system crashes during Analysis?  
During REDO?
- How do you limit the amount of work in REDO?
  - Flush asynchronously in the background.
  - Watch “hot spots”!
- How do you limit the amount of work in UNDO?
  - Avoid long-running Xacts.



## What should we learn today?



- Predict what portions of the log and database are necessary for recovery based on the recovery equations
- Explain how write-ahead logging is achieved in the ARIES protocol
- Explain the functions of recovery metadata such as the transaction table and the dirty page table
- Predict how recovery metadata is updated during normal operation
- Interpret the contents of the log resulting from ARIES normal operation
- Explain the three phases of ARIES crash recovery: analysis, redo, and undo
- Predict how recovery metadata, system state, and the log are updated during recovery

